

R2: Cross-Code Validation of the UH MkV FEL Beamline Optimisation

Eremey Valetov

4 March 2026

1 Introduction

The UH MkV FEL transport line (118 elements, 23 quadrupoles) was optimised independently using three simulation codes with distinct physics models: FELsim (linear transfer matrices), COSY INFINITY (differential algebra maps with fringe fields), and RF-Track (full particle tracking). Achieving agreement required identifying and correcting implementation errors in each code's adapter layer. This report documents the corrections applied, presents the cross-code comparison at $\varepsilon_n = 8 \pi \cdot \text{mm} \cdot \text{mrad}$, and benchmarks the optimisation algorithms.

The undulator Twiss targets follow Weinberg, Fisher & Li [1]: $\beta_x = 1.400 \text{ m}$, $\alpha_x = 0.470$, $\beta_y = 0.242 \text{ m}$, $\alpha_y = 0$. The 11-stage sequential optimisation decomposes the 23-quadrupole problem into physically motivated sections (doublets, triplets, chromaticity quads).

2 Simulation Codes

FELsim uses 4×4 linear transfer matrices with analytical dipole models: sector-bend bodies with a triangle-rule fringe correction for vertical edge kicks. The optimiser is Nelder-Mead (NM) with multi-start for Stage 11. $E = 40 \text{ MeV}$, 500 particles, seed = 42.

COSY INFINITY uses differential algebra (DA) to compute transfer maps to arbitrary order. The optimiser is the built-in gradient-based FIT (algorithm 1), which exploits DA derivatives. Two fringe field settings were tested: FR 0 (hard-edge, no fringe) and FR 1 (first-order Enge-model fringe, warm-started from FR 0). Computation order 3, 3 DA dimensions (6D phase space).

RF-Track (v2.5.5, CERN) performs element-by-element particle tracking through the beamline. Twiss parameters are computed from tracked particle distributions. The optimiser is NM with prefix caching: elements 0–87 are tracked once, and only the Stage 11 suffix (elements 87–118) is re-tracked per NM evaluation.

3 Model Corrections and Bug Fixes

Achieving cross-code agreement required identifying and correcting errors in each code's interface. Table 1 summarises the corrections.

Table 1: Bug fixes and model corrections enabling cross-code agreement.

Code	Issue	Fix
FELsim	DPW edge-kick sign inversion	$R = L/ \theta $ (unsigned)
RF-Track	SBend P/δ confusion	Analytical sector-bend correction
RF-Track	DPW modelled as Drift (L)	Thin-lens quad for edge kicks
RF-Track	Edge-kick sign (signed K_0)	$K_0 = \theta /L$
RF-Track	R_{12} unit conversion error	R-matrix invariant under scaling
RF-Track	Coord5 ct $\rightarrow$ l conversion (C7)	Corrected sign and scaling
COSY	Initial $\beta_0 = 1$ m default	$\beta_0 = \sigma^2/\varepsilon \approx 6.34$ m
COSY	Individual FIT objectives reset	Combined MSE objective
COSY	FR 1 cold-start failure	Warm-start from FR 0 solution

3.1 FELsim: DPW Edge-Kick Sign

FELsim’s `dipole_wedge` class computed the bending radius as $R = L/\theta$, using the signed bending angle. For negative-angle chicane dipoles ($\theta < 0$), this inverted R , flipping the sign of the vertical edge kicks. The fix replaces $R = L/\theta$ with $R = L/|\theta|$. All prior optimisation results involving chicane dipoles were invalidated by this bug.

3.2 RF-Track Adapter Corrections

Five issues were identified and resolved in the RF-Track adapter.

SBend P/δ confusion. RF-Track’s `SBend(L, angle, P_Q)` constructor interprets the bunch’s absolute momentum P (MeV/c, 6th column of `Bunch6d`) as the momentum deviation $\delta = \Delta P/P_0$, producing ~ 910 mm displacement. The setter-only `set_K0` has no effect. We implemented an analytical workaround: dipole bodies are tracked as `Drift(L)`, and the correction matrix $M_{\text{sector}} \times M_{\text{drift}}^{-1}$ is applied to (x, x') , with dispersion (R_{16} , R_{26}) and path-length (R_{56}) terms added. The correction is exact for the sector-bend transfer matrix (verified to 6 decimal places).

DPW modelling. The adapter converted dipole wedges to `Drift(L)` with the physical wedge length (≈ 10 mm), adding spurious position propagation. Replaced with thin-lens `Quadrupole(L=1e-10, K1L=-|K0|tan η)` for edge kicks.

Edge-kick sign. The `_annotate_dipole_edges()` method used signed $K_0 = \theta/L$, inverting edge kicks for negative-angle chicane dipoles. Fixed to $K_0 = |\theta|/L$.

R-matrix units. Erroneous $\times 1000$ scaling of R_{12} and $\div 1000$ scaling of R_{21} when converting between (m, rad) and (mm, mrad) coordinates. The R-matrix is invariant under this coordinate scaling because both position and angle scale by the same factor.

Coordinate conversion (C7). The FELsim $\leftrightarrow$ RF-Track coordinate conversion for the 5th phase-space coordinate (ct vs l) had incorrect sign and scaling, inflating the initial $\sigma(\text{ct})$ by $9.5\times$. This invalidated all longitudinal RF-Track results prior to the C7 fix.

3.3 COSY INFINITY Adapter Corrections

Initial Twiss. The initial beta function must be computed from the beam distribution: $\beta_0 = \sigma_x^2/\varepsilon \approx 6.34\text{ m}$ (where $\sigma_x = 0.8\text{ mm}$, $\varepsilon = 0.101\pi\text{ mm mrad}$). Using the common default $\beta_0 = 1\text{ m}$ produces completely incorrect transport ($\alpha_x = -2.8$ instead of ≈ 0 after Stage 1).

Combined MSE objective. COSY’s FIT with individual objectives (ENDFIT with multiple OBJ) resets all variables when *any* single objective fails to converge. A combined MSE objective preserves DA derivatives for gradient-based optimisation and converges monotonically.

Warm-starting for fringe fields. Running FR 1 from default initial guesses fails ($\text{MSE} \approx 0.2$) because the changed dipole edge kicks push the sequential FIT stages into incompatible local minima. Using the converged FR 0 currents as initial values for FR 1 resolves this: $\text{MSE} = 7.9 \times 10^{-8}$ with $\Delta I_{\text{max}} = 0.76\text{ A}$.

4 Cross-Code Baseline Results

Table 2 compares the optimised Twiss parameters at the undulator entrance for $\varepsilon_n = 8\pi\text{ mm mrad}$. All codes match the targets in the x -plane. The β_y residual in RF-Track (0.055 vs target 0.242 m) arises from the missing triangle-rule fringe correction ϕ in the thin-lens edge-kick model.

Table 2: Cross-code comparison at $\varepsilon_n = 8\pi\text{ mm mrad}$, $\sigma_E = 0.5\%$. Targets: $\beta_x = 1.400\text{ m}$, $\alpha_x = 0.470$, $\beta_y = 0.242\text{ m}$, $\alpha_y = 0$.

Code	Model	MSE	β_x (m)	α_x	β_y (m)	α_y
FELsim NM	Transfer matrix	1.20×10^{-6}	1.4000	0.4700	0.2418	0.0000
COSY FR 0	DA map (no fringe)	2.27×10^{-7}	1.3996	0.4700	0.2412	-0.0006
COSY FR 1	DA map (1st fringe)	3.65×10^{-9}	1.4000	0.4700	0.2418	0.0001
RF-Track opt	Particle tracking	7.02×10^{-3}	1.4010	0.4628	0.0548	-0.0028

FELsim NM and COSY (FR 0, FR 1) achieve MSEs of 10^{-6} – 10^{-9} , matching all four Twiss targets to within 10^{-3} . COSY’s gradient-based DA optimiser converges $\sim 10^3\times$ more precisely than NM. RF-Track’s higher MSE (7.0×10^{-3}) reflects the systematic β_y deficit from the thin-lens edge-kick model, not optimizer limitations.

Figure 2 shows the Twiss comparison graphically.

5 Quadrupole Current Comparison

Table 3 compares the optimised quadrupole currents across all four runs. Stages 1–2 and 6–9 show excellent agreement ($\Delta < 0.1\text{ A}$). The principal discrepancy is Stage 5 (double triplet): COSY finds a negative-polarity local minimum ($I_{35}, I_{37} < 0$), inaccessible to FELsim’s bounded NM ($I \in [0, 10]\text{ A}$). Stage 11 currents differ across all codes because the models produce different edge kicks in the final triplet section.

Despite the different Stage 11 currents, all codes match the undulator Twiss targets. The different optimal currents reflect genuine model differences — each code finds the best solution within its own physics. Direct current injection between codes (using one code’s optimised currents

Table 3: Optimised quadrupole currents: FELsim NM, COSY FR 0, COSY FR 1 (warm-started), and RF-Track opt (Stage 11 only).

Elem	Label	F/D	FELsim (A)	COSY FR 0 (A)	COSY FR 1 (A)	RFT-opt (A)	Stage
1	LQ1	F	0.822	0.839	0.842	—	1
3	LQ2	D	1.043	1.052	1.057	—	1
10	DPHQ	F	3.883	3.957	3.979	—	2
16	F1QA	D	2.240	1.993	1.994	—	3
18	F1QB	F	4.953	4.859	4.955	—	3
20	F1QC	D	3.426	3.499	3.599	—	3
27	DPHQ	F	4.666	4.685	4.711	—	4
33	F2QA	D	2.694	0.357	0.518	—	5
35	F2QB	F	2.652	−1.975	−1.849	—	5
37	F2QC	D	0.277	−1.995	−1.985	—	5
50	DPHQ	F	4.674	4.685	4.696	—	6
56	F3QA	D	3.122	3.168	3.218	—	7
58	F3QB	F	3.313	3.321	3.361	—	7
61	F3QC	F	5.178	5.166	5.252	—	8
63	F3QD	D	4.043	4.011	4.080	—	8
70	DPHQ	F	4.682	4.662	4.811	—	9
76	F4QA	D	3.934	3.869	3.953	—	10
78	F4QB	F	4.079	4.452	4.506	—	10
80	F4QC	D	0.014	0.608	0.592	—	10
87	DPHQ	F	1.362	1.259	1.610	2.045	11
93	UBQ1	D	0.945	2.348	2.473	1.405	11
95	UBQ2	F	2.885	3.802	3.602	0.722	11
97	UBQ3	D	2.192	4.747	3.995	0.111	11

in another) produces poor results (RFT-val MSE = 2.0 at $\varepsilon_n = 8$) because the accumulated edge-kick differences shift the beamline state at the undulator entrance.

6 Single-Dipole Transfer Matrix Comparison

Table 4 compares the vertical edge kick M_{43} for individual DPW-DPH-DPW dipole sandwiches. The discrepancy grows with bending angle: 3.9% at 1.5° , 8.9% at 4° , and 10% at 11.25° (FR 0 vs FELsim). The FC Enge coefficients for the FC1/FC2 chicane produce 76% weaker y -kicks, indicating that realistic fringe fields partially cancel the geometric edge focusing.

Table 4: Vertical edge kick M_{43} (rad/m) and Y-plane stability for representative dipole triplets.

Dipole	θ ($^\circ$)	FELsim	COSY FR 0	COSY FR 3	COSY FR 3+FC
M_{43} (vertical kick, rad/m)					
Transport	1.5	-7.42×10^{-3}	-7.71×10^{-3}	-6.78×10^{-3}	—
Chicane	4.0	-1.11×10^{-1}	-1.21×10^{-1}	-8.93×10^{-2}	—
FC1	11.25	-9.50×10^{-1}	−1.04	-7.79×10^{-1}	-2.27×10^{-1}
$ \text{Tr}(M_y)/2 $ (stability, < 1 required)					
Transport	1.5	0.9997	0.9997	0.9997	—
Chicane	4.0	0.9977	0.9975	0.9982	—
FC1	11.25	0.9823	0.9805	0.9855	0.9957

7 Parameter Sensitivity

Table 5 summarises the S4 parameter scans (energy spread, chirp, emittance) using the MSE quality thresholds: Excellent ($< 10^{-3}$), Acceptable (< 0.01), Marginal (< 0.1), Failed (≥ 0.1). The optimisation is robust across typical operating ranges. The only failure occurs at extreme emittance values ($\varepsilon_n \leq 2$ or $\varepsilon_n = 5$, where NM is trapped in a local minimum).

Table 5: Quality distribution across S4 parameter scans (FELsim NM).

Scan	Range	N	Exc.	Acc.	Mar.	Fail	Best MSE	Worst MSE
σ_E (%)	0.1–5	15	14	1	0	0	8.52×10^{-7}	1.08×10^{-3}
h (s $^{-1}$)	0–4e10	12	12	0	0	0	1.32×10^{-6}	2.28×10^{-4}
ε_n (mm mrad)	1–20	10	6	3	0	1	1.20×10^{-6}	2.39×10^1

Figure 1 shows the MSE vs parameter for all three scans.

8 Optimiser Comparison

Two evolutionary optimiser configurations (glyfada) were benchmarked against NM for Stage 11 at three emittance points. Table 6 compares the MSE values. COSY FIT results are included as a reference for the best achievable MSE at each emittance.

8.1 W6: Glyfada ULS (Uniform Random)

The initial benchmark (W6) used glyfada’s default ULS algorithm with uniform random initialisation over the full parameter space ($[0, 15] \times [0, 10]^3$), $30 \times 20 = 600$ evaluations. NM outperforms by 3–6 orders of magnitude at all emittance points. At $\varepsilon_n = 5$, all 600 evaluations return the penalty value (unstable optics).

8.2 W7: Glyfada CMA-ES (Warm-Started)

A re-benchmark (W7) used CMA-ES with warm-starting from the NM solution, informed bounds (± 3 A around NM), feasibility-rules constraint handling, and $5 \times$ more evaluations (600 total). Despite these improvements, CMA-ES still fails catastrophically: $\text{MSE} = 7.7 \times 10^5$ at $\varepsilon_n = 5$ and $\text{MSE} = 17.8$ at $\varepsilon_n = 8$.

Table 6: Stage 11 optimiser comparison across emittance points. COSY FIT uses gradient-based DA optimisation; FELsim NM uses Nelder-Mead with 5 restarts; Glyfada ULS and CMA-ES use evolutionary search.

ε_n	COSY FIT	FELsim NM	Glyfada ULS	Glyfada CMA-ES
5	1.40×10^{-2}	4.02×10^{-1}	1.00×10^6	7.73×10^5
8	2.27×10^{-7}	1.20×10^{-6}	2.96	1.78×10^1
14	1.44×10^{-8}	1.09×10^{-3}	3.11×10^{-1}	2.56

The FELsim MSE landscape has an extremely narrow feasibility basin: large regions of the 4D parameter space produce divergent optics. Evolutionary search cannot navigate this landscape

efficiently, even with warm-starting and tight bounds. The gradient-based COSY FIT (exploiting DA derivatives) is the most effective optimiser, followed by NM with multi-start.

Figure 4 shows the MSE comparison graphically.

9 Summary

1. **Three-code agreement achieved.** FELsim, COSY INFINITY, and RF-Track independently match the undulator Twiss targets at $\varepsilon_n = 8$, with MSEs ranging from 3.7×10^{-9} (COSY FR 1) to 7.0×10^{-3} (RF-Track). Table 2 and Fig. 2.
2. **Nine bugs fixed across three codes** (Table 1). The most critical were the FELsim DPW edge-kick sign, the RF-Track SBend P/δ confusion, and the COSY initial Twiss computation. All prior chicane-related optimisation results were invalidated.
3. **Different quad currents, same Twiss.** Each code converges to a different set of Stage 11 currents because the physics models produce different dipole edge kicks (Table 3). Direct current injection between codes fails; each code must optimise within its own model.
4. **Dipole model is the dominant source of discrepancy.** Single-dipole comparisons show 4–20% differences in the vertical edge kick depending on bending angle and fringe field model (Table 4).
5. **Optimisation is robust** across energy spread (0.1–5%) and chirp ($0\text{--}4 \times 10^{10} \text{ s}^{-1}$), with $\text{MSE} < 10^{-3}$ everywhere except extreme emittance values (Table 5).
6. **Gradient-based optimisation (COSY FIT) outperforms all others.** Evolutionary search (glyfada) fails at every emittance point by 3–6 orders of magnitude vs NM, even with CMA-ES, warm-starting, and tight bounds (Table 6). The narrow feasibility basin of the FELsim MSE landscape makes evolutionary approaches ineffective for this problem.

Figures

References

- [1] M. Weinberg, A. Fisher, and R. Li, “Design of the UH Mark V Free Electron Laser,” arXiv:2510.14061v1, Table I.




Figure 1: MSE vs parameter for the S4 sensitivity scans: energy spread (left), chirp (centre), emittance (right). The dashed lines indicate quality thresholds.




Figure 2: Cross-code Twiss comparison at $\epsilon_n = 8$. All codes match β_x and α_x closely. RF-Track's β_y deficit (0.055 vs 0.242 m) is due to the missing triangle-rule fringe correction in the thin-lens edge-kick model.




Figure 3: Bunch compression factor vs chirp rate (W12 Part A). Analytical and COSY map propagation agree closely. The compression floor (≈ 0.45 ps) prevents reaching 0.5 ps at practical chirp values.




Figure 4: Stage 11 MSE comparison across optimisers and emittance points (W6/W7 data). COSY FIT (gradient-based DA) achieves the lowest MSE at $\varepsilon_n = 8$ and 14; NM is competitive. Both glyfada configurations fail by orders of magnitude.